

Name:

email adres:

Student id:

Educational Program:

Exercise 3: measurements of the diameter of arteries in angiograms

We develop and evaluate image processing for measuring the diameters of a blood vessel in coronary angiographic images such as the one shown in Figure 3. The exercise consists of two parts:

- part I: edge detection and morphological operation
Hand-on training for multi-scale edge detection, such as Canny and Marr-Hildreth.
- part II: assessment of the accuracy of the diameter measurements
Edge detection is useful for delineation of objects. But it has also limitations, which we will examine.



Figure 1. An angiogram showing the image of artery

Part I: Edge based image segmentation

1. The scale space

A 2D Gaussian point spread function can be created and visualized with:

```
sigma = 5;                                % set the width of the Gaussian
L = 2*ceil(sigma*..)+1;                   % fill in a constant to define the matrix size
xymax = (L-1)/2;                          % the maximum of the x and the y coordinate
xrange = -xymax:xymax;                   % the range of x values
yrange = xrange;                         % the range of y values
% create the PSF matrix
h = fspecial('gaussian', L, sigma);

%% visualize as a 3D plot:

% create a RGB matrix to define the colour of the surface plot
C = cat(3, ones(size(h)), ones(size(h)), zeros(size(h)));

% create the surface plot of the gaussian
figure(1)
hd = surf(xrange, yrange, h, C, 'FaceColor', 'interp', 'Facelight', 'phong');
camlight right                            % add a light at the right side of the scene
xlim([...]);                             % set appropriate axis limits
ylim([...]);
xlabel('x');                              % add axis labels
ylabel('y');
zlabel('h(x,y)');
print -r150 -dpng ex3_1.png % print the result to file
```

The function `fspecial` is used to create a Gaussian PSF matrix `h` with the deviation (width) of the Gaussian set to σ (matlab: `sigma`). The size of `h` should be such that the truncation error of the tails of the Gaussian is negligible. On the other hand, a too large size leads to a waste of computations. The function `surf` creates a surface plot of the PSF. The 3D array `C` defines the colour of the surface. The other functions, `camlight`, `xlabel`, etc, are needed to refine your graph¹.

- a) Initiate a new matlab script. That is, start with `%% ex 3 - name`, then clear the work space and close all figure windows, then add a line “ `%% question 1`. Next copy and paste the code given above². Complete the unfinished lines.

Give the completed Matlab code of the 2nd line above:

¹ **Note:** in the sequel, omission of axis labels in any graph (images excluded) in your report will be considered as a shortcoming of your report, and will have a negative impact on the grade.

² Maybe you have to retype part of the text as the adobe pdf may use other fonts than that the Matlab editor expects.

Insert the graph:




- b) Apply the filter to the test image stored in `ang2.png`. Apply this with four different values of σ . That is: $\sigma = 1, 10, 20$ and 35 . You can use the following template for your code:

```
im = ...                % read the image from file
sigma = [1 2 4 8];      % fill array with sigma values
for i=1:4                % do 4 times:
    h = fspecial(... ;    % create the PSF
    imfiltered = imfilter(... ; % apply the filter
    subplot(2,2,i);        % define a subplot
    imshow(imfiltered, []); % show the image
    title(['\sigma = ' num2str(sigma(i))]); % include a plot title
end
```

Select the options of `imfilter` such that the output image has the same size as the input image, and the boundary option is such that boundary effects on the border of the image are visually minimized.

Give the completed Matlab code of the lines that contain the `fspecial` function and `imfilter` function:

Insert the resulting images:



What are the scales of these four images?

2. Derivatives

The first derivative of a continuous image $f(x, y)$ is denoted by $f_x(x, y) = \partial f(x, y) / \partial x$. In digital image processing, only digital representations $f(n\Delta, m\Delta)$ and $f_x(n\Delta, m\Delta)$ are available. The derivative $f_x(n\Delta, m\Delta)$ is only an approximation. In the discrete case, differentiation is necessarily extended to a finite neighbourhood. It can be approximated by a convolution that implements the differentiation combined with some sort of low-pass filtering.

In the continuous domain this combination is nicely accomplished by:

$$\frac{\partial}{\partial x} (f(x, y) * h(x, y)) = f(x, y) * \left(\frac{\partial}{\partial x} h(x, y) \right)$$

With that, differentiation in the digital domain is accomplished by:

$$f_x(n\Delta, m\Delta) \equiv f(n\Delta, m\Delta) * h_x(n\Delta, m\Delta)$$

In the scale space theory it has been argued that the best low-pass filter needed for differentiation is the Gaussian:

$$h(x, y) = \frac{1}{2\pi\sigma^2} \exp\left(-\frac{x^2 + y^2}{2\sigma^2}\right) \quad (1)$$

Code to create the PSF matrix h with a size of $L \times L$ (as an alternative to using the function `fspecial`) is:

```
N = ceil((L-1)/2); % get the size of half of the full range
[x,y]=meshgrid(-N:N,-N:N); % create the coordinates of a 2D orthogonal grid
h = exp(-(x.^2 + y.^2)/(2*sigma^2))/(2*pi*sigma^2);
```

- a) Give the analytical expression for $h_y(x, y)$. That is, differentiate expression (1) analytically with respect to y . Insert the found expression in Matlab syntax such as in the sample code for $h(x, y)$:

- b) Create a PSF matrix h_y that contains a sampled version $h_y(n\Delta, m\Delta)$ of $h_y(x, y)$. Use $\sigma = 1.5\Delta$ (which is equivalent to $\sigma = 1.5$ and $\Delta = 1$). Visualize the PSF as in question 1, and insert the figure. Apply the convolution with h_y to the artery image, insert the resulting image. Use the function `imwrite` and `mat2gray`, rather than a copy of a figure window)

- c) Explain the response of the operator to:

- Areas without much contrast.
- Horizontal step transitions such as at the boundaries of vertical blood vessels.
- Vertical step transitions such as at the boundaries of horizontal blood vessels.

INTERMEZZO

The usage of the matlab function `mat2gray`:

Remember:

- If the number representation of an image is of type `uint8`:
 - The range of pixel is: 0, 1, 2, ..., 255. So, the resolution of the grey scale is 1.
 - 0 → black and 255 → white
- If the number representation is `double`:
 - The range of pixels is between 0 and 1, and the resolution of the grey scale is very fine (10^{-16})
 - 0 → black and 1 → white

When writing an image `im` to file (jpg or tif), then:

- If the number representation of `im` is `uint8`, the pixels are written to file without processing. So, the bytes of the image are literally written to the file (apart from a possible compression).
- If the number representation of `im` is `double`, then the pixels are rescaled by a factor of 255, that is: `im` becomes `im*255`, then converted to `uint8`, and finally written to file.


In the latter case, information may be lost. First of all, there will be round-off errors. Secondly, if the grey level of a pixel was larger than 1 (whiter than white), or smaller than 0 (blackier than black), then in the conversion from `double` to `uint8`, the pixel will be first truncated to 1 or to 0, respectively.

This truncation can be prevented by the function `mat2gray`. This function maps the greyscale of an image linearly (`im` becomes `a*im+b`) such that the maximum of `im` becomes 1 and the minimum becomes 0. This guarantees that no truncation takes place when writing to a file, but it also may enlarge the round-off error.


**FROM NOW IT IS UNDERSTOOD THAT YOU USE MAT2GRAY WHENEVER NEEDED !
FAILING TO DO SO HAS A NEGATIVE IMPACT ON THE GRADE**

3. The function `ut_gauss` is an easy-to-use implementation³ to calculate the derivatives $f_x(n, m)$, $f_y(n, m)$, $f_{xx}(n, m)$, $f_{yy}(n, m)$, and $f_{xy}(n, m)$. Apply these operations to the test image. Use $\sigma = 3$.


fx:



fy:



fxx:



fyy:



fxy:



³ We silently assume that $\Delta = 1$ from now on.

Consider a vertical artery. Which of these images convey information about this artery, and what kind of information is that? Which of these images do not convey information in this particular case?



4. Gradient magnitude and Laplacian

The gradient magnitude of an image is defined as:

$$|\nabla f(n, m)| = \sqrt{f_x^2(n, m) + f_y^2(n, m)}$$

or, in short, $\sqrt{f_x^2 + f_y^2}$. The Laplacian is defined as:


$$\Delta f(n, m) = f_{xx}(n, m) + f_{yy}(n, m)$$

or shortly $f_{xx} + f_{yy}$.


Using the results from question 3, calculate the gradient magnitude and the Laplacian of the test image.

Don't use for-loops.


gradient magnitude



Laplacian:



How can these images help to find edge segments?



5. Marr-Hildreth's zero crossings

In the continuous domain the equation

$$f_{xx}(x, y) + f_{yy}(x, y) = 0$$

defines, under mild conditions, curves in the x, y -plane. These curves are called the *zero crossings* of the Laplacian.


The goal of the Marr-Hildreth operation is to find an edge map. The edge map is a binary image in which all found edge elements are marked by 1 where all other pixels are 0. Marr and Hildreth mark each pixel as an edge element if the Laplacian exhibits a zero crossing near that pixel position.

A cheap zero crossing detection is accomplished by first finding areas with a positive Laplacian, and then to find the boundary of these areas. We apply this procedure to the Laplacian of the test image obtained in question 4:


- a) Create a binary image in which a pixel is set to 1 if the corresponding Laplacian for that pixel is positive. **(Don't use for-loops!)**
- b) Determine the boundary pixels of the binary image, and display it on the screen. For that purpose:
 - Create a structuring element consisting of a 4-neighbourhood (use `strel` with option `diamond`.)
 - Erode the binary image by this structuring element, so that only the interior of the foreground is left.
 - Subtract this interior from the original binary image.

Note: `bwmorph` with option `'remove'` does almost the same job but is still less useful as it operates differently on the border of the image).

thresholded Laplacian:




zero crossings Laplacian:




- c) Form the resulting zero crossings 8-connected paths or 4-connected paths?
- d) Instead of marking areas with positive Laplacian values, areas with negative Laplacians can be marked. Repeat a) and b) with these negative areas.

thresholded Laplacian:



zero crossings Laplacian:



- e) If you compare the results of question 5b with the one of 5a, which of the two will underestimate the diameter of an artery, and which of the two will overestimate it?

- f) To prevent the under- or overestimation, mentioned in question e, the two results will be combined by means of a logical or operation. In matlab, this is the operation `|`. Calculate this, and show the results below. The next step is to skeletonize the found structure to assure paths with a thickness of 1 pixel.

Combined zeros crossings:

Skeletonized:

6. Hysteresis thresholding the gradient magnitude

The zero crossings of the Marr-Hildreth operator produces many candidate edge elements, but many of them will be false. The reason for this is that the Laplacian may exhibit a zero crossing at positions for which the gradient magnitude is near zero. So, the condition for a zero crossing to be an acceptable edge element is that the gradient magnitude exceeds some suitable threshold.

- Determine the gradient magnitude as in question 4. Mask this image by the Marr-Hildreth zero crossings obtained in question 5. That is, all gradient magnitude pixels that are not a zero crossing should be made zero, while all other pixels are retained.
- Compare the masked gradient magnitude image against a suitable threshold T ; and show an edge map consisting of zero crossings for which the gradient magnitude is larger than T . Choose the threshold such that the narrowing in the centre is just visible in the edge map. Show the result on the screen. You can interactively select the threshold by visually evaluating the edge map.


Found threshold:

Insert found edge map:

In hysteresis thresholding, the gradient magnitude is first masked by the zero crossings (so as to exclude all pixels which are not a zero crossing, as before). Next, the result is thresholded twice resulting in two edge maps. The first map, called the `marker`, is created with a large threshold. It may miss many weak edges, but there will be only a few false edges. The second map, called the `mask`, is created with a lower threshold. It contains both weak and strong edges, but also many false edges. The strategy is to accept an edge segment in the second image only if such an edge segment contains at least one edge element from the first edge map. Thus, the `marker` will be propagated into the `mask` to yield a more reliable edge map.

- c) Apply a threshold operation with a threshold that is somewhat larger than the previously found threshold T . Name the resulting image `marker`.
- d) Apply a threshold operation with a threshold T . Name the resulting image `mask`.
- e) Propagate the `marker` image into the `mask` image, and interpret the result. **Hint:** use the function `imreconstruct`.

Resulting edge map:




Describe the (potential) advantages of hysteresis thresholding:

7. Edge detection


The Marr-Hildreth, including hysteresis thresholding, is implemented in the function `ut_edge`. This function uses the function `ut_levelx` to find the zero crossings. `ut_levelx` implements a more accurate method to find the zero crossings than the one outlined in question 5. `ut_edge` also contains an implementation of the Canny edge detector.

- b) Apply `ut_edge` to the test image in order to get an edge map in which the edge elements form two non-fragmented edge segments. Select the options of `ut_edge` as follows:
 - Set the option 'c' That is, apply Canny. Later we will also use 'm' (Marr-Hildreth)
 - Use option 's' to set the scale. You can try it first with $\sigma = 4$. But this may not be the best choice.
 - Use the option 'h' (hysteresis thresholding). The thresholds are specified as two additional parameters. Read the help of the function to understand how these parameters are defined. Select these parameters and the scale interactively such that the narrowing is well delineated, but without too much spurious edges in the remaining part of the map. Perhaps, you need some iterations to establish the right parameters and scale.
 - Repeat with the Marr-Hildreth operator.

Canny edge map:



Marr Hildreth edge map:



- c) Describe in one sentence the main difference between the localization criterion of the Marr-Hildreth edge detector and the one of the Canny edge detector. Hint: read the powerpoint sheets.

- d) What is the theoretical advantage of the Canny's criterion above Marr-Hildreth criterion?

Part II: Accuracy of edge-based diameter measurements

In this part, an experiment will be conducted to quantitatively assess the quality of diameter estimation from X-ray images using edge detection. Software is available that simulates the imaging of an artery with a given diameter. An example of a test image is provided in Figure 2.




Figure 2. Test image based on a simulation of X-ray imaging of an artery

The test image shows a horizontally aligned artery. Prior knowledge of this alignment facilitates the determination of the diameters because only two horizontal edge segments needs to be localized. The assumption of horizontally aligned arteries is not very unrealistic as a suitable geometric transform can be applied that ensures this. See Figure 3 for an example.

The simulation of the test image is based on an accurate model of the imaging device. It includes:

- A circular model of the cross section of the blood vessel.
- The pixel size is $\Delta = 40 \mu\text{m}$. Pixels are square. A diameter of 1 mm corresponds to 25 pixels.
- The maximum exposure is equivalent to an average of 12500 quant/mm², which is 20 quant/pixel.
- The OTF of the imaging device, a so-called image intensifier, is given by:

$$H(\rho) = \exp\left(-\frac{\rho}{\rho_0}\right) \text{ with } \rho_0 = 2 \text{ lp/mm}$$

The Matlab function that accomplishes this simulation is `im_bloodvessel`. This function produces a 250x1000 image. The diameter of the simulated artery is defined in units of mm, and should be passed to the function as an input argument of the function. In each call of the function, a new random realisation of the noise is generated.




Figure 3 De-skewing a coronary angiogram

a) Original X-ray image with superimposed the centre line (green) of a blood vessel, and with a grid (red) of which the rows are orthogonally aligned across this centre line. The columns are aligned along the centre line.

b) The de-skewed image. The grid in a) is rearranged in an orthogonal grid. This aligns the artery horizontally in the centre of the image.

A second function, `get_diameters`, is available that calculates for each column the diameter, expressed in pixels. The help of this function is as follows:

```
[diameters,success]=get_diameters(im,sigma,op)
```

Estimation of the diameter of a horizontally aligned blood vessel

INPUT:

```
im:          input intensity image
sigma:       gaussian width of the applied filters in the edge detector
op:         edge operator. 'm' : marrhildreth, 'c' : canny
```

OUTPUT:

```
diameters    array with for each column in the image, the estimated diameter
success:     flag to indicate successful estimation. 'true' (1): success
```


Example:

```
im = im_bloodvessel(2);          % simulation of an artery with 2 mm diameters
[diam,suc] = get_diameters(im,3,'m'); % sigma=3 performed with Marr-Hildreth operator
```

If the flag `success` is false (0), the estimation was not successful, and the resulting diameters are not valid.

8. Generate a new test image with a diameter of 2 mm, and estimate with `get_diameters` the diameters using the canny operator and with $\sigma = 1$. Inspect and describe what happens with the success flag and the resulting diameters. Explain how these results should be interpreted in this particular case..

9. Estimate the diameters again, but now with $\sigma = 5$. Check the value of the success flag. Show the found diameters by plotting them in a graph. Don't forget the axis labels. Insert the graph:



The quality of the edge-based diameter estimation has four aspects:

- **Robustness**: the ability to produce results that make sense. Robustness can be expressed in the *success rate*: the probability of not having an outlying result caused by some internal fatal error.
- **Precision**: the ability to reproduce the same result if the measurement is repeated under the same conditions. This can be quantified by means of the *standard deviation* of the results.
- **Accuracy**: the ability of having a low systematic error. A systematic error is an error that replicates each time the experiment is repeated under the same conditions. This can be quantified by the *bias* of the results.
- **Detailedness**: the ability to resolve the curvedness of the boundaries of the artery.

This exercise will only address the first three aspects.

10. Calculate the standard deviation (function `std`) and the average of the diameters (function `mean`) of the estimated diameters in question 9.

Standard deviation of the estimated diameters:

Average diameter:

What is in this case the real diameter if expressed in pixels?

The bias of `get_estimators` can be roughly assessed by the difference between the average diameter and the real diameter. What is this difference (in pixels)?

Why is this only a rough assessment?

11. A failure of the estimator is internally detected inside `get_diameters`, and indicated by the flag `success`. To determine the success rate, a single image does not suffice. To assess the robustness, i.e. to estimate the probability of success, a number of images with different noise realizations are needed.

- a) Suppose that we have N images with different noise realizations, and suppose that from these N images `get_diameters` reports n images with successful results, give the expression for estimating the success rate (use Matlab syntax):

- b) If it is required, that the uncertainty of the estimated success rate is about 10%, how many images with different noise realizations would be approximately enough to achieve this uncertainty?

- c) Write Matlab code, in which the N images are generated, and in which consecutively the success rate is estimated. Do this with the following parameters: real diameter = 1.5 mm, canny operator, $\sigma = 4$. What is the resulting success rate?

12. In question 10, you calculated the standard deviation and bias of the estimator based on just one image. In question 11, a couple of images were generated to estimate the success rate. Now that you have these

images, the estimation of the standard deviation and bias can also be based on these multiple images, so as to improve the uncertainties of these estimates. Expand the code in question 11 to implement this.

Tips: you can do so by first collecting the results of `get_diameters` in N dimensional arrays, and then processing these arrays to get the estimates of the standard deviation and bias. Note that when `get_diameters` indicates a failure by means of the `success` flag, the values of the diameters are not valid.


13. The estimated success rate, standard deviation, and bias depends on:

- the real diameter D , expressed in mm.
 - the edge operator: canny 'c', or marr-hildreth 'm'.
 - the width of the edge operator σ .
- a) Set up an experiment, in which the success rate, standard deviation, and bias are estimated with varying width: $\sigma = 1, 2, \dots, 40$. Do this with $D = 1.5$ mm, and the canny operator. Repeat this with the marr-hildreth operator. Put the results of both operators in three graphs: 'success rate', 'standard deviation', and 'bias'. Don't forget the axes labels.
- b) Repeat this with $D = 3$ mm. Add the results to the 3 graphs. Don't forget to add a legend (**legend**).

Insert plot success rate :



Insert plot standard deviation:



Insert plot bias:



14. Observe typical and a typical behaviour in the resulting graphs. If possible, explain (put question mark if you can't explain):

a) Success rate:

Behaviour for small σ :

Behaviour for large σ :

Explanation behaviour with respect to σ :

Behaviour comparing canny and marr-hildreth::

Explanation behaviour with respect to canny and marr-hildreth:

Behaviour comparing $D = 1.5$ mm and $D = 3$ mm

Explanation behaviour with respect to diameters:

b) Standard deviation:

Behaviour for small σ :

Behaviour for large σ :

Explanation behaviour with respect to σ :

Behaviour comparing canny and marr-hildreth::

Explanation behaviour with respect to canny and marr-hildreth:

Behaviour comparing $D = 1.5$ mm and $D = 3$ mm

Explanation behaviour with respect to diameters:

c) Bias:

Behaviour for small σ :

Behaviour for large σ :

Explanation behaviour with respect to σ :

Behaviour comparing canny and marr-hildreth::

Explanation behaviour with respect to canny and marr-hildreth:

Behaviour comparing $D = 1.5$ mm and $D = 3$ mm

Explanation behaviour with respect to diameters:

15. What is/are the disadvantage(s) of using edge detection in this application?

Insert m-code (note a copy-and- paste of this code to Matlab’s native editor, and a single ‘run’ should suffice to run the whole code):